

The Simunto logo consists of the word "Simunto" in a white, rounded, sans-serif font, centered within a dark purple rounded rectangle.

Simunto

The DSim logo features the text "DSim" in a bold, teal-colored, sans-serif font.

**DSim**

A Distributed Message-Passing Mobsim Implementation

Janek Laudan, Christian Rakow, Marcel Rieser · Simunto GmbH  
Steffen Axer, Hannes Rewald · Volkswagen AG

# **Problem & Recap**

# What is our problem?

MATSim modelers want to:

- Simulate larger model domains
- With greater detail

This leads to increased computational demands and long runtimes as:

- We can only run on a single machine
- The QSim's current parallelization architecture does not take full advantage of modern multicore CPU hardware

# What could be the solution?

We presented: *Distributed parallel Qsim* (with **Paul Heinrich** and **Kai Nagel**)

- A prototype implementation of a distributed mobsim in Rust
- MPI based distributed computation
- Aimed at High-Performance Computing Hardware
- Scalability up to 1000 CPU-Cores
- [Open-Source Implementation](#)

Published as *High-Performance Mobility Simulation: Implementation of a Parallel Distributed Message-Passing Algorithm for MATSim*<sup>1</sup>

**Why DSIm?**

# Why DSim?

Implement a message-passing  
mobility simulation in Java

1. Take advantage of modern  
multicore CPUs
2. Stay compatible with the  
existing MATSim-Codebase
3. Support distributed execution  
of MATSim-Simulations

# Why DSim – General Idea

Take the main ideas from Rust-Prototype

- Adopt a message-passing architecture for the mobsim
- Distribute simulation work by partitioning the simulated network

While staying compatible to the ecosystem

- The simulation must execute in the JVM
- Current Infrastructure should be compatible
- The new mobsim implementation should be usable as a drop-in replacement

**Architecture**




# Architecture

Distribute Workload geographically

Domain Decomposition is a well-understood problem with plenty of algorithms, e.g., METIS

Estimate computational load for each vertice of the graph by estimating #vehicles crossing the node

METIS balances the computational load across partitions




# Architecture

Distribute Workload geographically



Domain Decomposition is a well-understood problem with plenty of algorithms, e.g., METIS

Estimate computational load for each vertice of the graph by estimating #vehicles crossing the node

METIS balances the computational load across partitions




# Architecture



# Architecture

One MATSim process is started per JVM

- Multiple simulation processes
- Multiple events handlers
- One Communicator
- One Events Manager
- One Task Scheduler




# Architecture

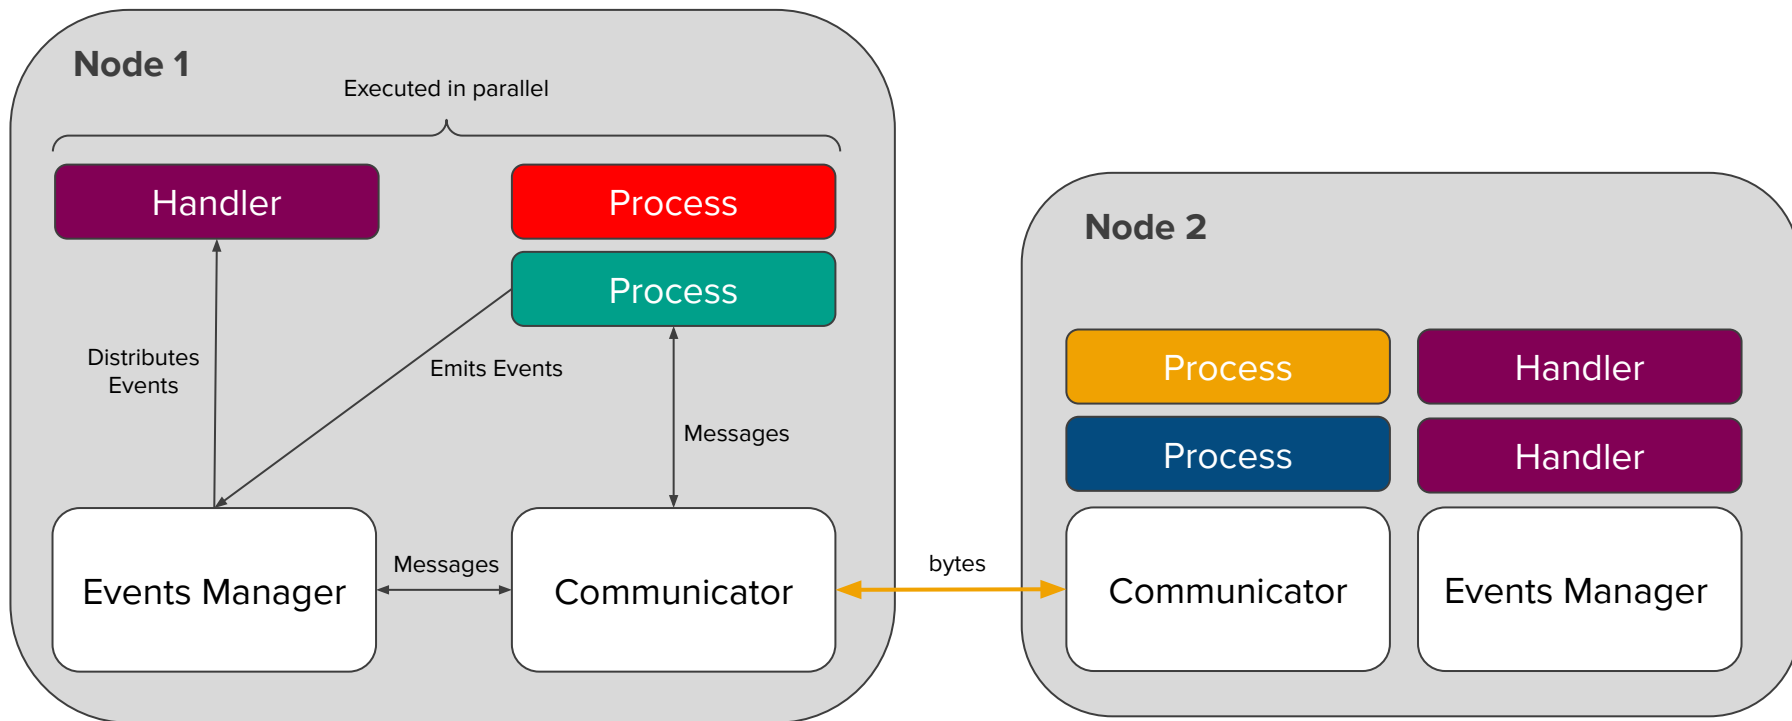
One MATSim process is started per JVM

- Multiple simulation processes
- Multiple events handlers
- One Communicator
- One Events Manager
- One Task Scheduler

Multiple processes communicate via the network layer



# Architecture



# Architecture – Gotchas

## Unified Parallelization Strategy

- Mobsim and event handling are part of the same executor
- Parallelization on the ‘outer’ loop
- Simulation process is single threaded
- Event Handlers are single threaded
- Well defined synchronization via messages

Possibility for distributed traffic simulation

Possibility for distributed event processing

# Results



# Benchmark Scenario


Kanton Zürich with 1.6 million synthetic  
persons



Demand Generation with  
Creario

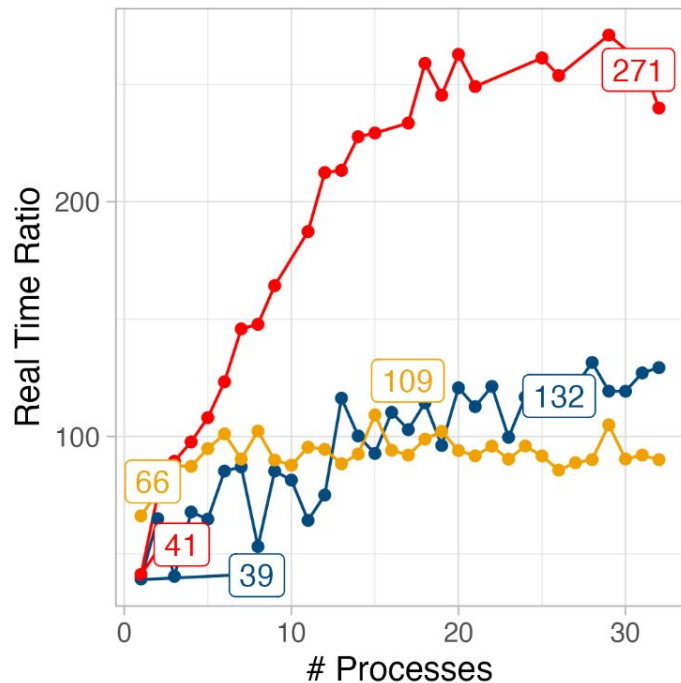
Partitioning with METIS based on  
expected node weights

Test Environment: High-Performance  
Cluster with AMD EPYC™ 7313 (2x16  
cores) and 240GB RAM

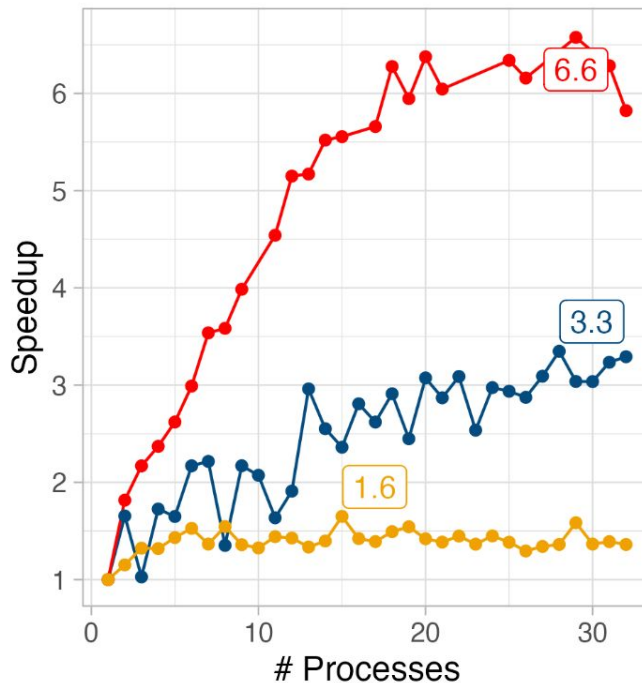


# Benchmark Kanton Zürich

## Real Time Ratio



## Relative Speedup

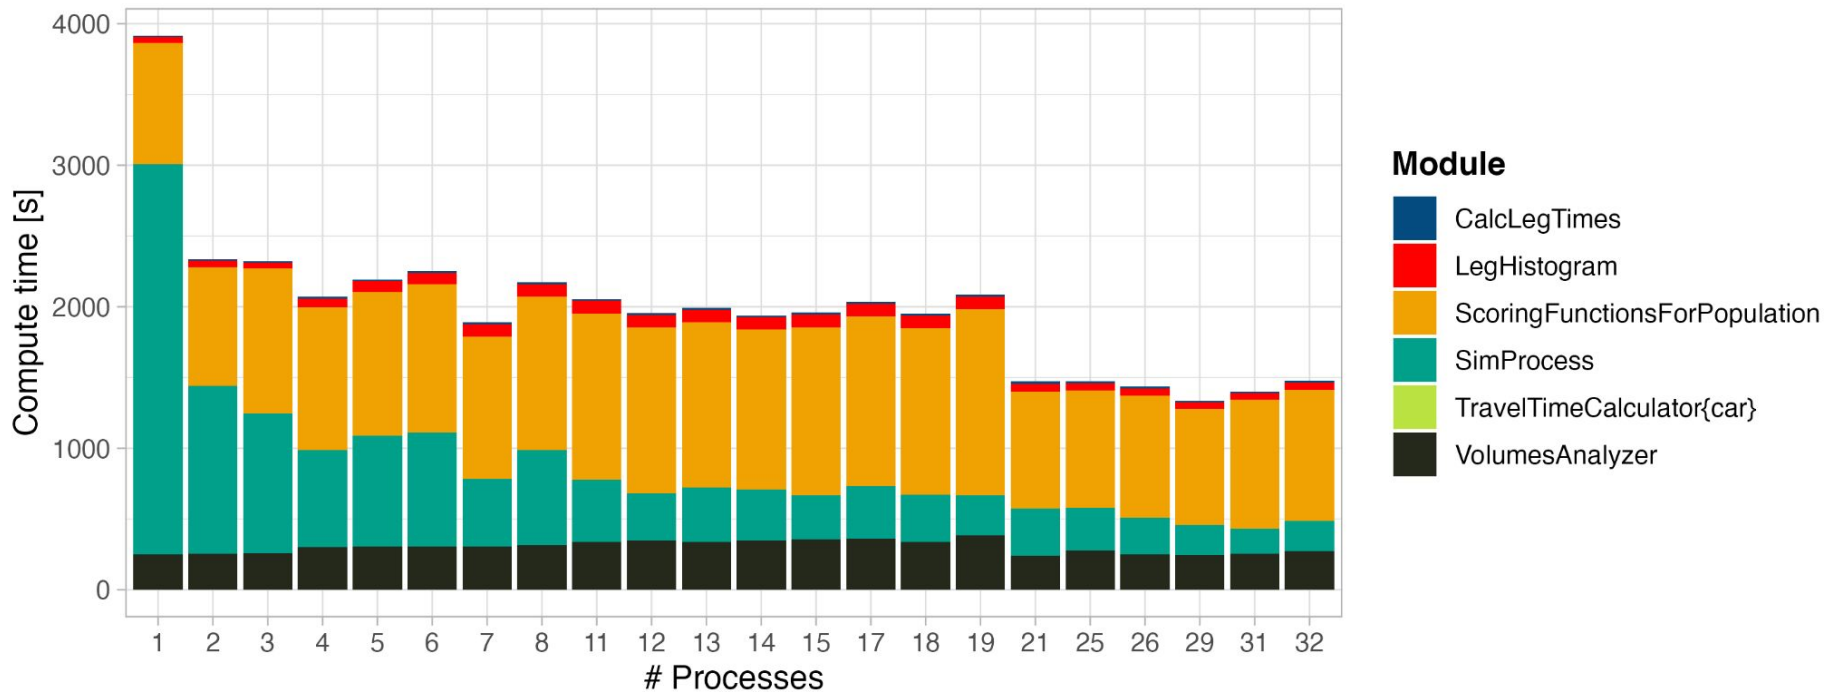


### Setup

- Dsim -- Unweighted Partitioning
- Dsim -- Weighted Partitioning
- QSim


# Benchmark Kanton Zürich

## Summed durations on Partition 0




# Benchmark DRT (Berlin Scenrio)

## Real Time Ratio with DRT



## Speedup with DRT



### Setup

- DSIm DRT RSI
- DSIm DRT SI
- QSim DRT RSI
- QSim DRT SI

# Conclusion

1. Architecture for distributed computing in MATSim implemented
  2. Improved scalability compared to QSim
  3. Better reasoning about bottlenecks
  4. Truly distributed execution needs further testing and improved stability
-

# Outlook

Help and Support appreciated

Adapt more parts of MATSim to take advantage of new architecture

- Event handling must be switched to a distributed approach
- Develop strategy for services such as DRT

DSim will be merged into the main repository in the upcoming weeks

---

# Closing

Janek Laudan

Software Consultant @ Simunto

[laudan@simunto.com](mailto:laudan@simunto.com)

<https://simunto.com>

Simunto GmbH

Riedgrabenweg 49

8050 Zürich

Schweiz

[mail@simunto.com](mailto:mail@simunto.com)

Thanks to Steffen Axer's group at Volkswagen  
for supporting this work financially!

